

Retrieval-Augmented Generation (RAG)

RAG combines LLMs with external data retrieval.

Key components: Retrieval, Augmentation, Generation.

LangChain supports RAG via Document Loaders (like PyPDFLoader).